Larger-network inference with newly trained models

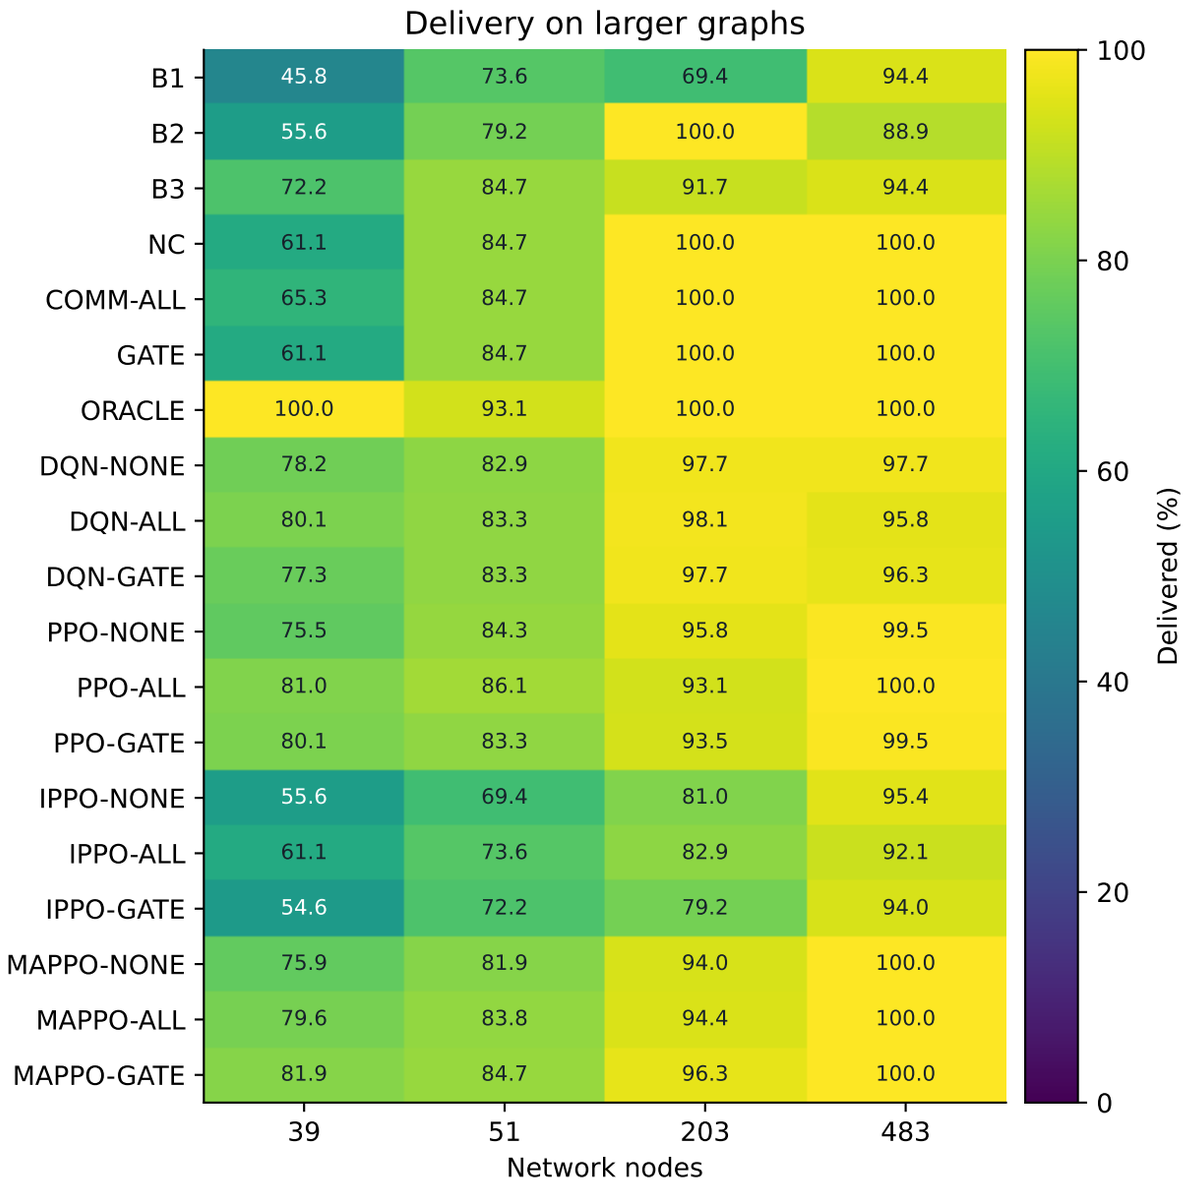

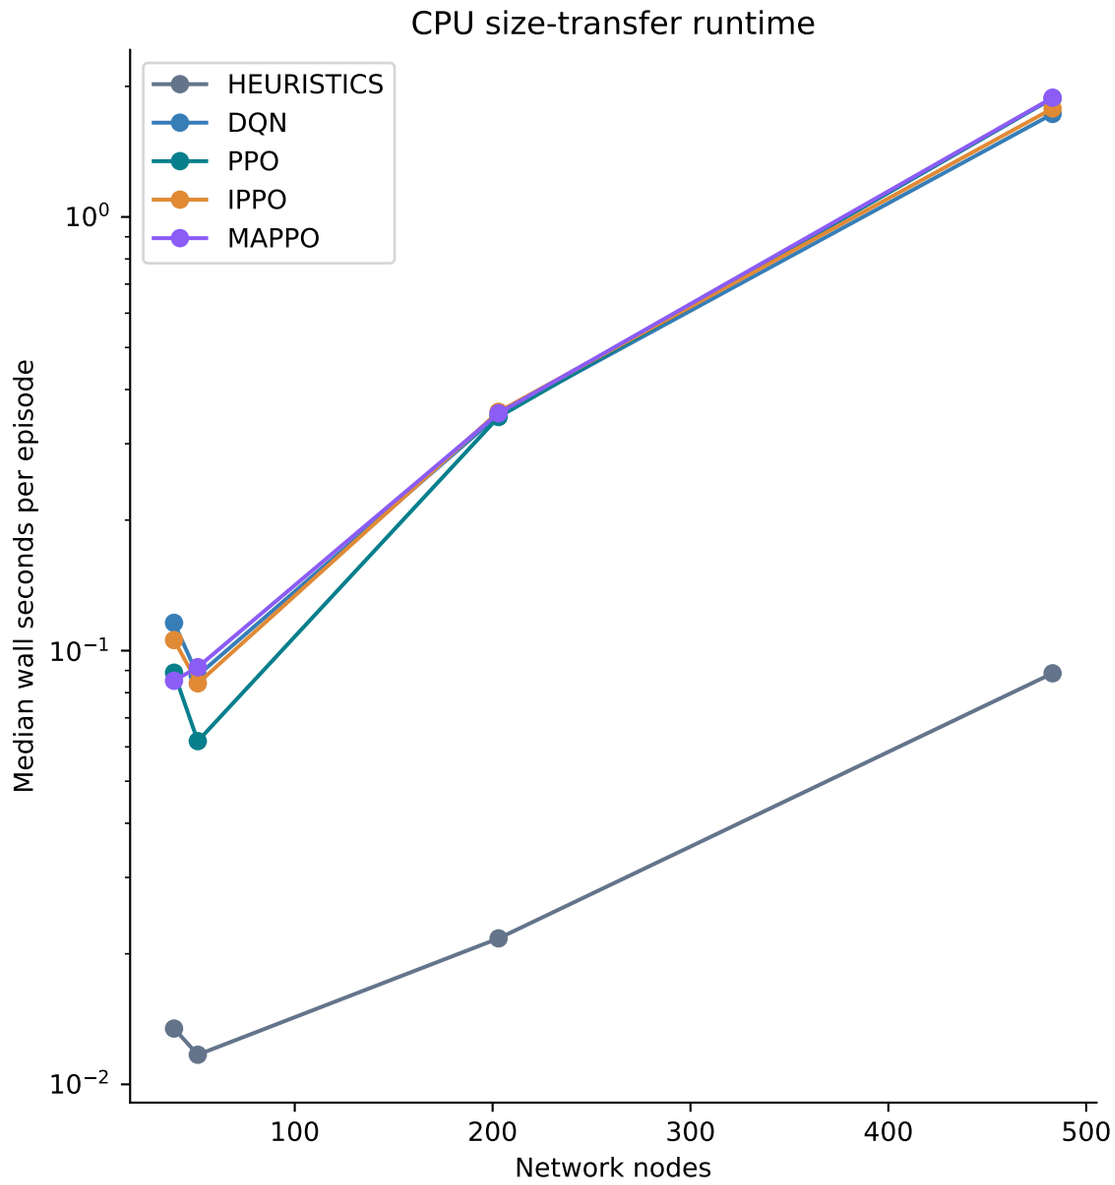

3 graphs × 2 episodes per size; population/attack budget fixed. Threat density therefore varies. Timing excludes Python allocation instrumentation.

EXPLORATORY RESULTS • 300 episodes/model • 3 training seeds • convergence not assumed