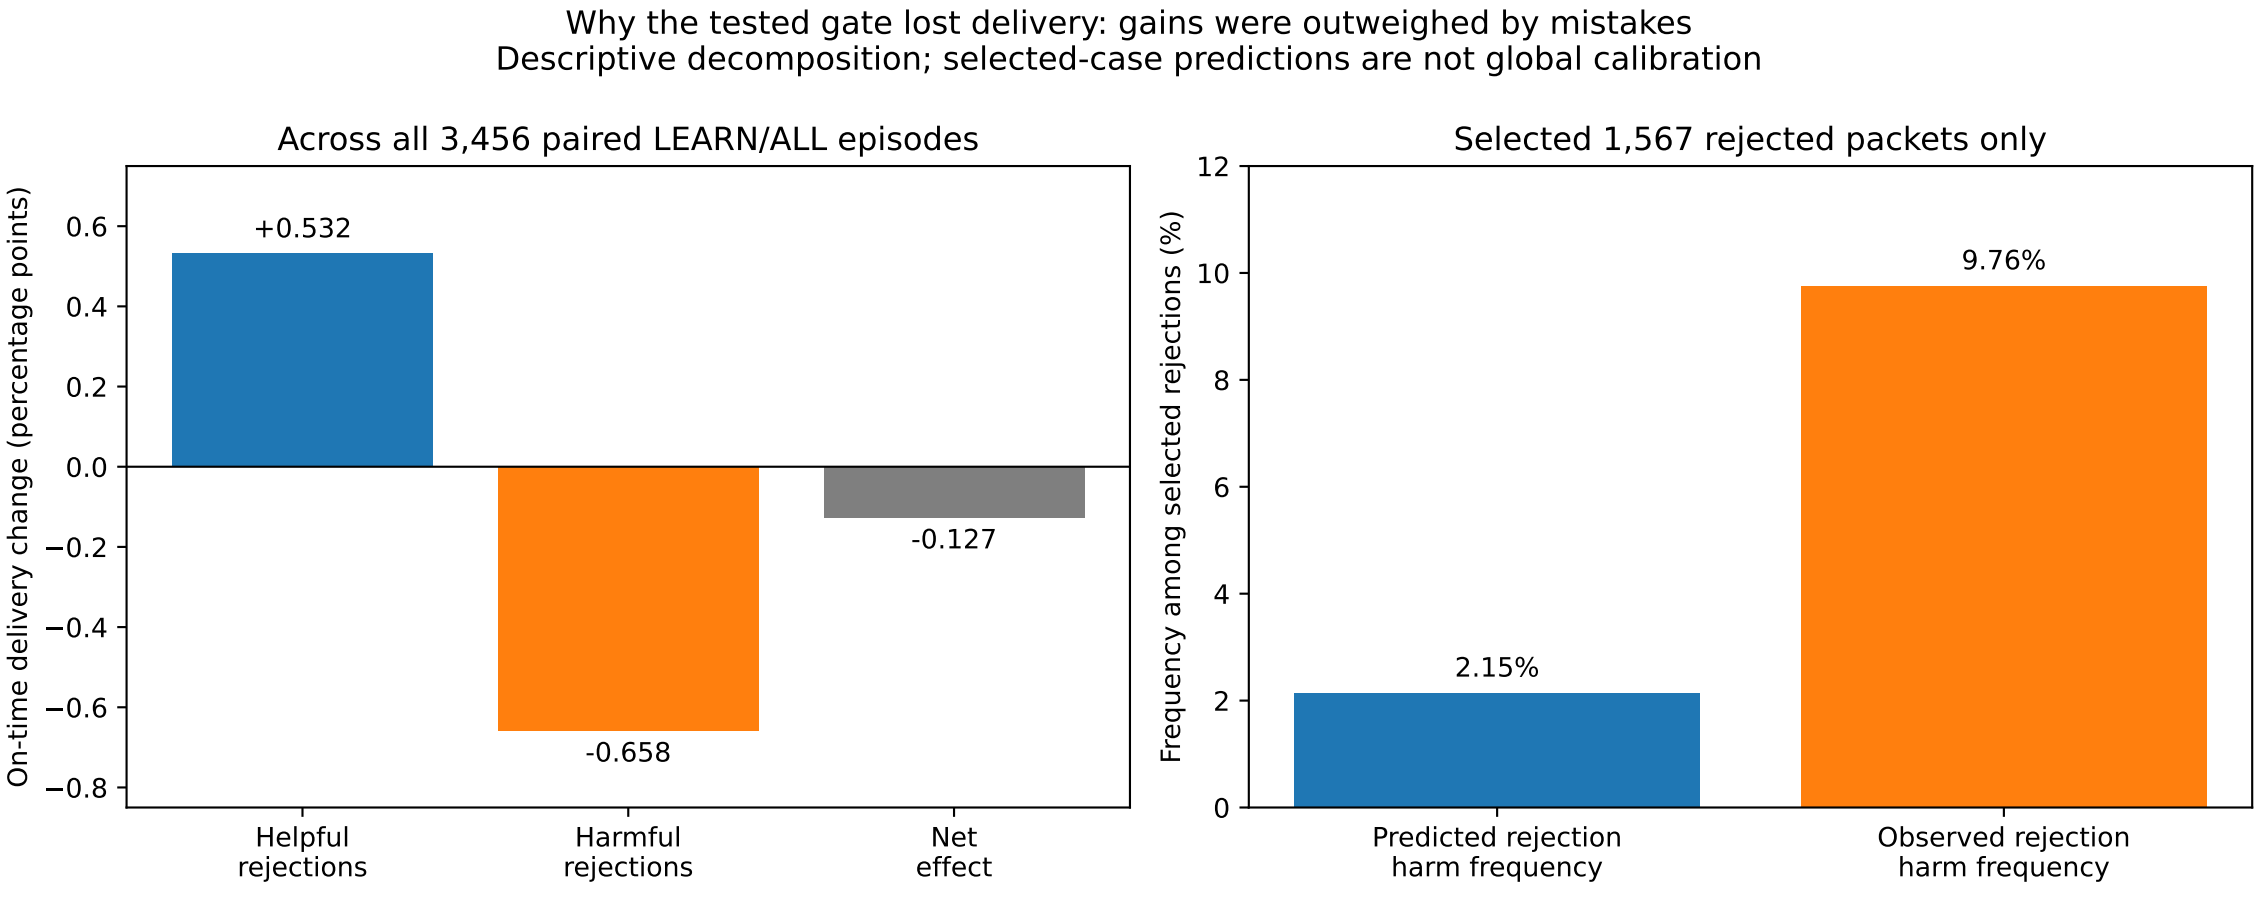
Why the tested gate lost delivery: gains were outweighed by mistakes
Descriptive decomposition; selected-case predictions are not global calibration
Across all 3,456 paired LEARN/ALL episodes
On-time delivery change (percentage points)
+0.532
-0.658
-0.127
Helpful rejections
Harmful rejections
Net effect
Selected 1,567 rejected packets only
Frequency among selected rejections (%)
2.15%
9.76%
Predicted rejection harm frequency
Observed rejection harm frequency