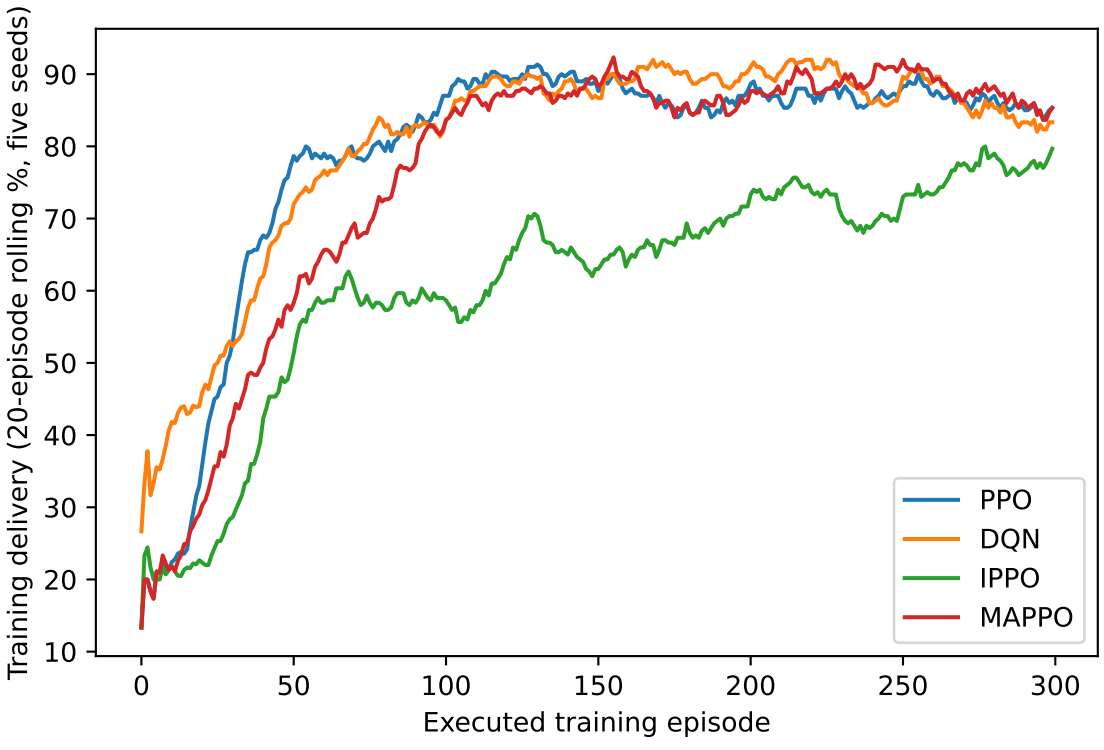
Training delivery (20-episode rolling %, five seeds)
10
20
30
40
50
60
70
80
90
0
50
100
150
200
250
300
Executed training episode
PPO
DQN
IPPO
MAPPO

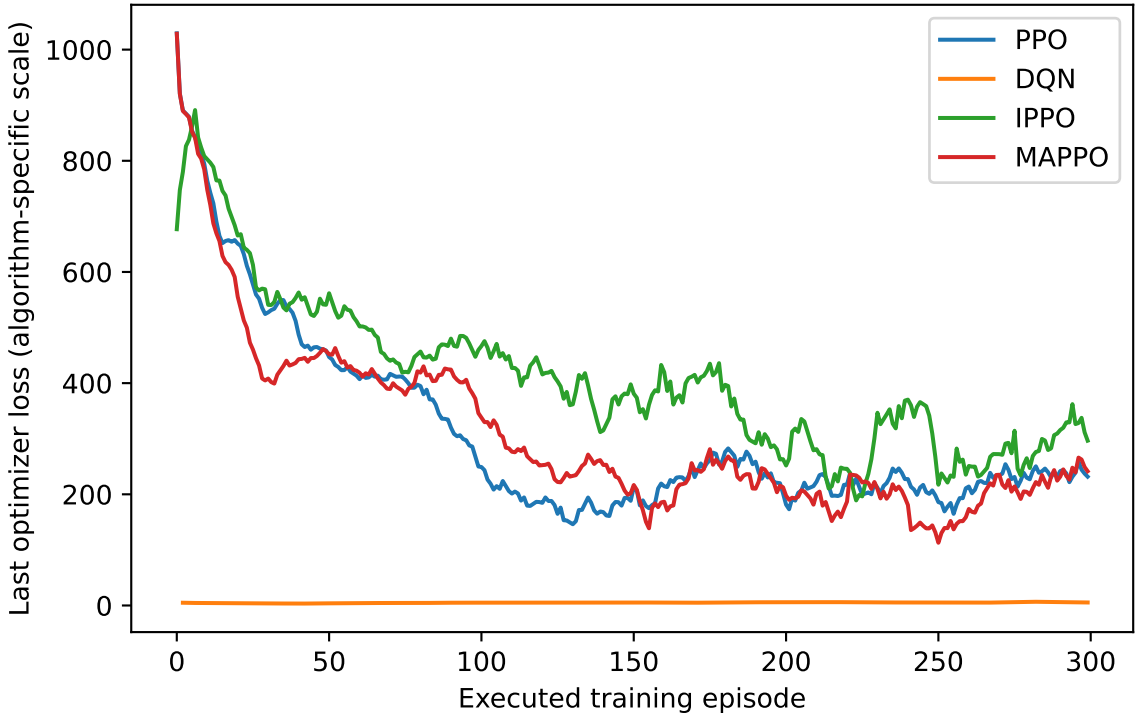
Last optimizer loss (algorithm-specific scale)
0
200
400
600
800
1000
0
50
100
150
200
250
300
Executed training episode
PPO
DQN
IPPO
MAPPO