Longer training under contested conditions

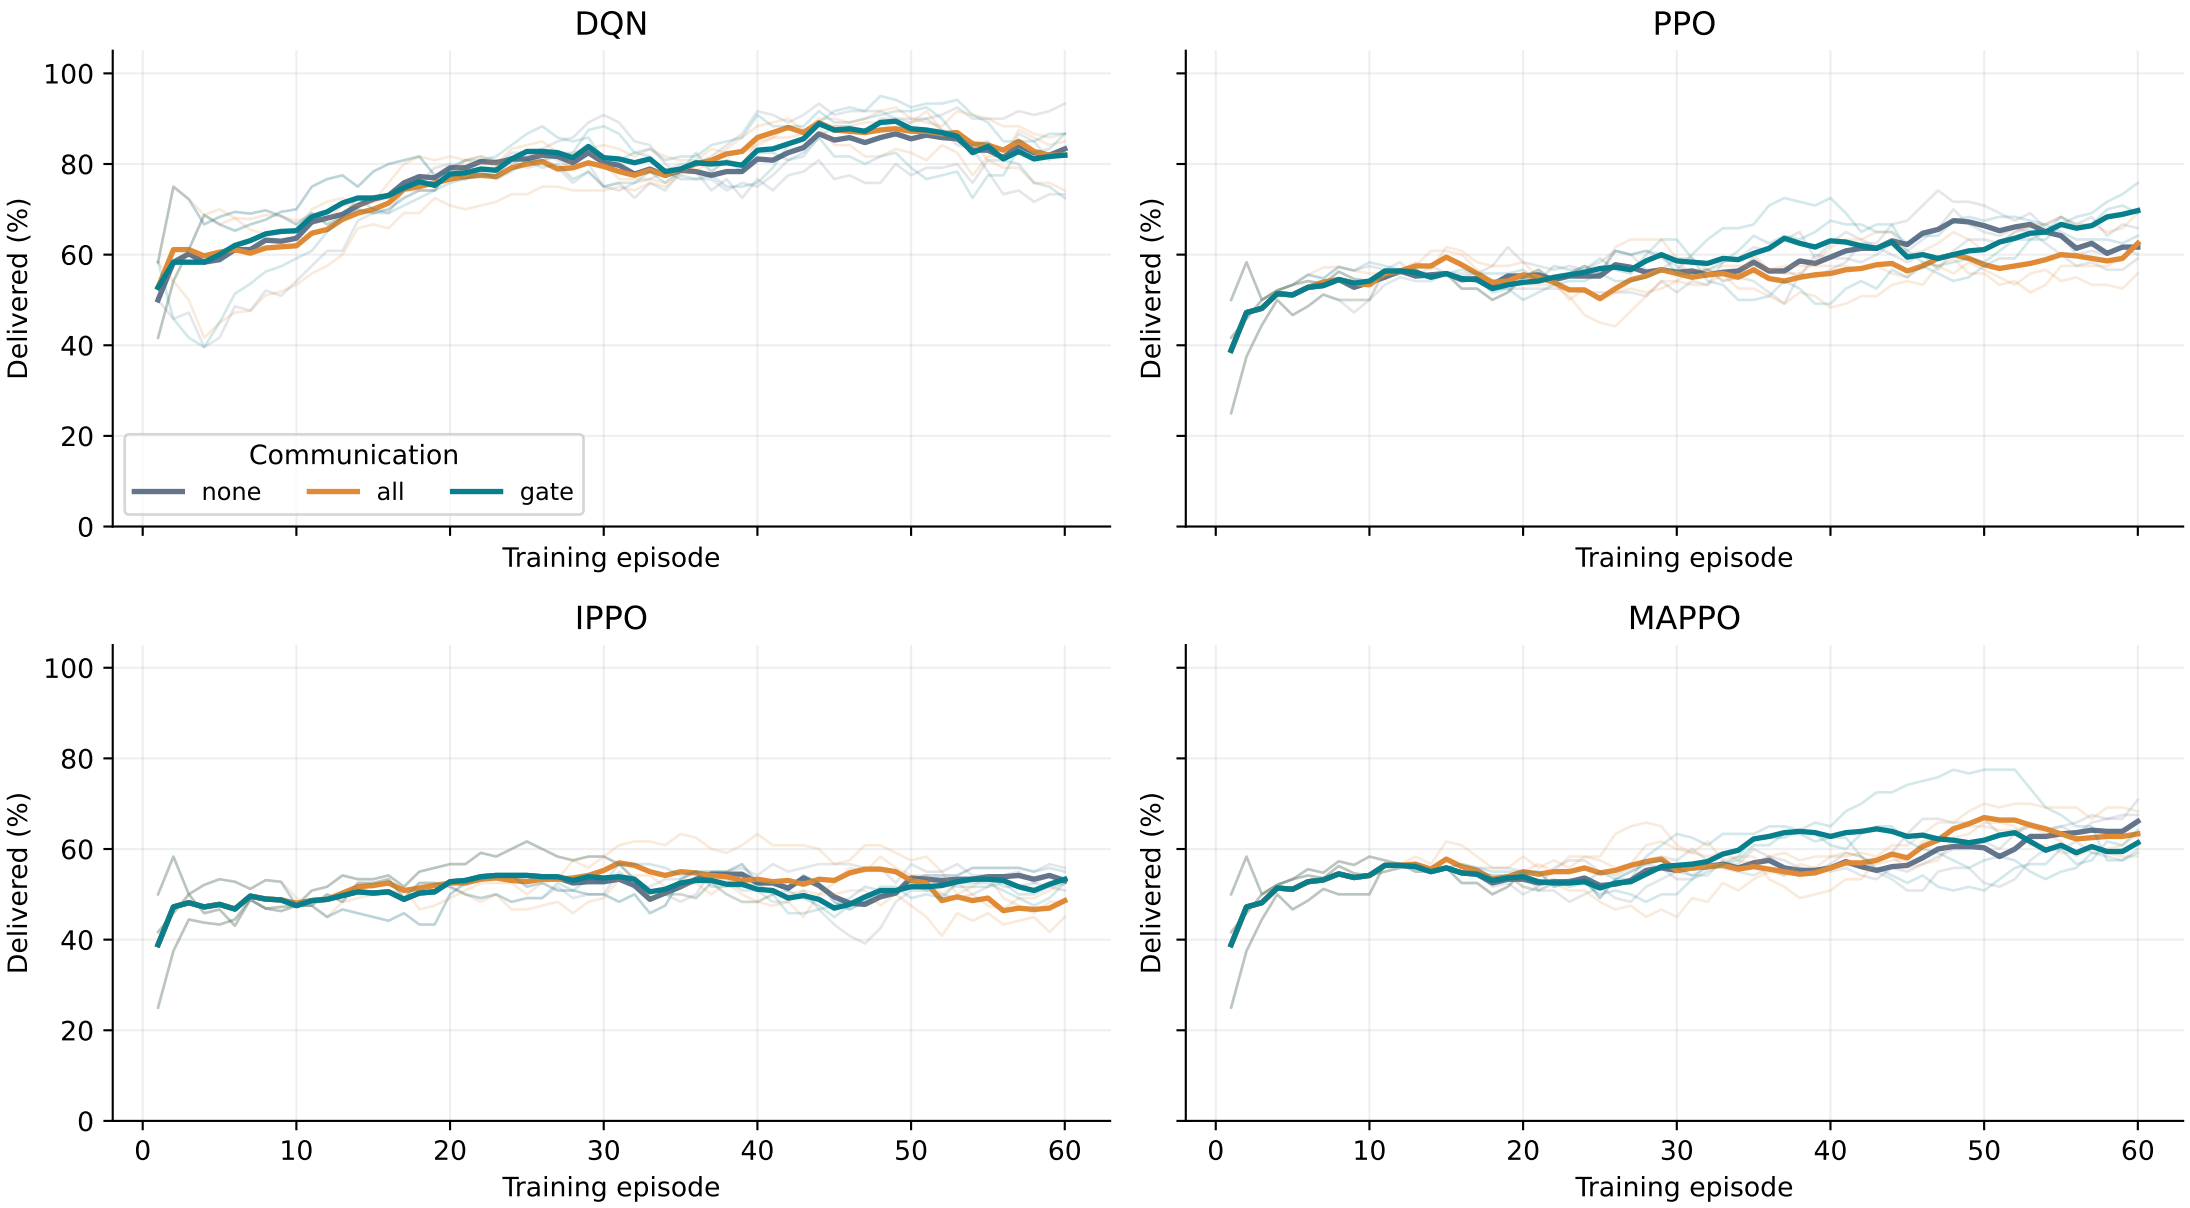

10-episode moving averages. Bold: three-seed mean; faint: individual seeds. Stochastic training outcomes, not fixed-policy convergence tests.

EXPLORATORY RERUN • 60 episodes/model • 3 training seeds • convergence not assumed