Matched held-out comparisons with uncertainty

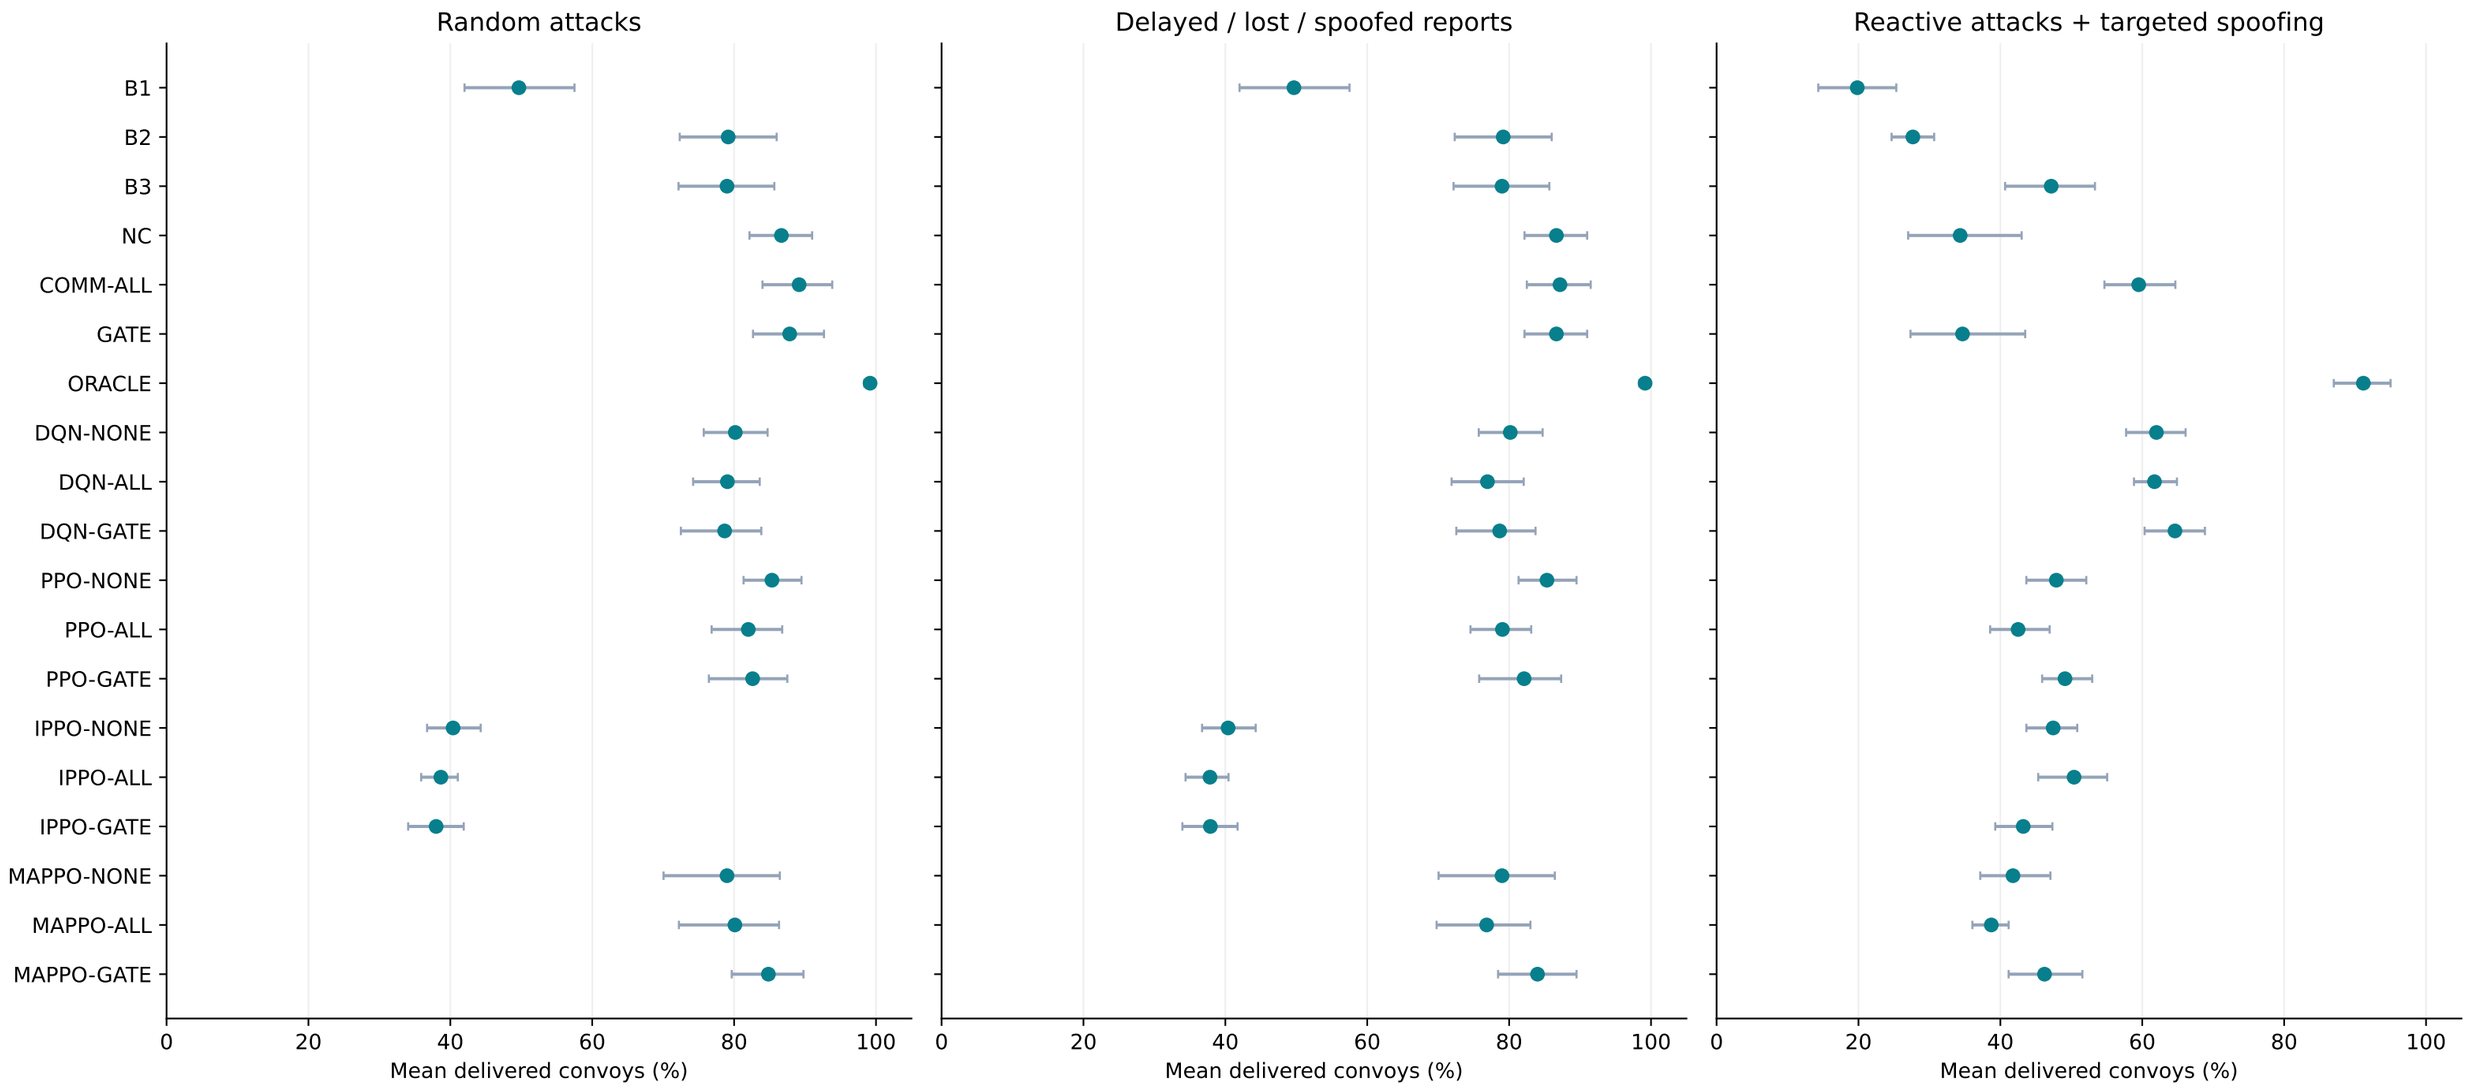

Same 10 held-out graphs × 5 episodes across scenarios. Bars: graph-cluster 95% bootstrap intervals conditional on trained models; 12 convoys.

EXPLORATORY RERUN • 60 episodes/model • 3 training seeds • convergence not assumed